

Bike-Sharing Business Analysis (Python Version)

Date: September 18, Thursday to September 23, Tuesday, 2025

Aims & Goals: This .ipynb file serves as a supplement to the Bike-Sharing Business Analysis project, showcasing a fundamental analytical workflow using Python. Additionally, a comparative analysis between using Python and Google Sheets for this bike-sharing business analysis is conducted during the project.

Self-directed professional: Qi Zhou

Email: qiqizhou1996@gmail.com

LinkedIn: www.linkedin.com/in/qi-zhou-1996to2096

Github: <https://github.com/QiZhou1996/a-bike-sharing-business-analysis-project>

Introduction:

This Bike-Sharing Business Analysis (Python Version) consists of three sequential steps: first, importing packages and loading the dataset; second, understanding the dataset through descriptive statistics and data cleaning; and lastly, performing exploratory data analysis (EDA) and data visualizations. For the insight report of this analysis, please see the above *LinkedIn* or *Github* link.

From data cleaning, including removing missing values, identifying duplicates, and handling outliers, to visualizations, comprising boxplots, bar charts, pie charts, countplots and scatterplots, key findings are summarized and compared (with those in Google Sheets).

Data Analysis using Python:

Step 1: Import Packages and Load Dataset

```
In [1]: # Import packages

# For data manipulation
import numpy as np
import pandas as pd

# For data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Load dataset into a dataframe
raw_data = pd.read_csv("raw-data.csv")

# Display the first five rows of the raw data and check if the data is lo
raw_data.head()
```

Out [1]:

	Start Time	Stop Time	Start Station ID	Start Station Name	End Station ID	End Station Name	Bike ID	User Type	Birth Year	Age
0	01-17 00:38	1-1-17 01:03	3194	McGinley Square	3271	Danforth Light Rail	24668	Subscriber	1961	6
1	01-17 01:47	01-17 01:58	3183	Exchange Place	3203	Hamilton Park	26167	Subscriber	1993	2
2	01-17 01:47	01-17 01:58	3183	Exchange Place	3203	Hamilton Park	26167	Subscriber	1993	2
3	01-17 01:56	01-17 02:00	3186	Grove St PATH	3270	Jersey & 6th St	24604	Subscriber	1970	5
4	1-1-17 02:12	01-17 02:23	3270	Jersey & 6th St	3206	Hilltop	24641	Subscriber	1978	4

Step 2: Understand the Dataset through Descriptive Statistics and Data Cleaning

In [2]: `# Gather the basic information about the raw data`
`raw_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20400 entries, 0 to 20399
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Start Time                            20400 non-null  object
1   Stop Time                             20400 non-null  object
2   Start Station ID                      20400 non-null  int64
3   Start Station Name                    20400 non-null  object
4   End Station ID                        20400 non-null  int64
5   End Station Name                      20399 non-null  object
6   Bike ID                              20400 non-null  int64
7   User Type                            20400 non-null  object
8   Birth Year                            20400 non-null  int64
9   Age                                  20400 non-null  int64
10  Age Groups                            20400 non-null  object
11  Trip Duration                         20400 non-null  int64
12  Trip_Duration_in_min                  20400 non-null  object
13  Month                                 20400 non-null  int64
14  Season                               20400 non-null  object
15  Temperature                           20400 non-null  int64
16  Weekday                              20400 non-null  object
dtypes: int64(8), object(9)
memory usage: 2.6+ MB
```

In [3]: `# Gather the descriptive statistics about the raw data`
`raw_data.describe(include='all')`

Out [3]:

	Start Time	Stop Time	Start Station ID	Start Station Name	End Station ID	End Station Name	Bike
count	20400	20400	20400.000000	20400	20400.000000	20399	20400.000000
unique	15746	16039	NaN	50	NaN	56	NaN
top	20-03-17 17:39	09-03-17 08:20	NaN	Grove St PATH	NaN	Grove St PATH	NaN
freq	7	7	NaN	2544	NaN	3313	NaN
mean	NaN	NaN	3215.863627	NaN	3211.439510	NaN	25301.7326
std	NaN	NaN	34.563120	NaN	82.707121	NaN	989.9742
min	NaN	NaN	3183.000000	NaN	152.000000	NaN	15084.000000
25%	NaN	NaN	3186.000000	NaN	3186.000000	NaN	24523.000000
50%	NaN	NaN	3203.000000	NaN	3202.000000	NaN	24679.000000
75%	NaN	NaN	3267.000000	NaN	3220.000000	NaN	26220.000000
max	NaN	NaN	3281.000000	NaN	3442.000000	NaN	29296.000000

In [4]: *# Check missing values*
raw_data.isna().sum()

Out [4]:

Start Time	0
Stop Time	0
Start Station ID	0
Start Station Name	0
End Station ID	0
End Station Name	1
Bike ID	0
User Type	0
Birth Year	0
Age	0
Age Groups	0
Trip Duration	0
Trip_Duration_in_min	0
Month	0
Season	0
Temperature	0
Weekday	0
dtype: int64	

Finding 1:

There is one missing value that can be removed, matching the way missing values are handled in Google Sheets.

In [5]: *# Remove missing values*
raw_dropna = raw_data.dropna()
#raw_dropna.info()

In [6]: *# Check duplicates*
raw_dropna.duplicated().sum()

```
# Check some rows containing duplicates as needed
#raw_dropna[raw_dropna.duplicated()].head()
```

```
Out [6]: np.int64(1950)
```

Finding 2:

Note that Pandas here identifies 1,950 duplicates, whereas Google Sheets identifies 3,555 duplicates.

On the one hand, it is critically important to pay attention to whether different methods of handling duplicates affect the analysis results; on the other hand, from the perspective of de-duplicating techniques, it is worth further study.

```
In [7]: # Remove duplicates
data_deduplicated = raw_dropna.drop_duplicates(keep='first')
#data_deduplicated.info()
```

```
In [21]: # Check outliers of the "Trip_Duration_in_min" variable
data_deduplicated.loc[:, 'Trip Duration (min)'] = data_deduplicated.loc[:,
data_deduplicated.info()
sns.boxplot(x=data_deduplicated['Trip Duration (min)'])
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 18449 entries, 0 to 20399
```


```
Data columns (total 18 columns):
```

#	Column	Non-Null Count	Dtype
0	Start Time	18449 non-null	object
1	Stop Time	18449 non-null	object
2	Start Station ID	18449 non-null	int64
3	Start Station Name	18449 non-null	object
4	End Station ID	18449 non-null	int64
5	End Station Name	18449 non-null	object
6	Bike ID	18449 non-null	int64
7	User Type	18449 non-null	object
8	Birth Year	18449 non-null	int64
9	Age	18449 non-null	int64
10	Age Groups	18449 non-null	object
11	Trip Duration	18449 non-null	int64
12	Trip_Duration_in_min	18449 non-null	object
13	Month	18449 non-null	int64
14	Season	18449 non-null	object
15	Temperature	18449 non-null	int64
16	Weekday	18449 non-null	object
17	Trip Duration (min)	18449 non-null	int64

```
dtypes: int64(9), object(9)
```

```
memory usage: 2.7+ MB
```

```
Out[21]: <Axes: xlabel='Trip Duration (min)'>
```



Finding 3:

Note that converting the *Object* variable of "Trip_Duration_in_min" to the *Int* variable of "Trip Duration (min)" leads to no box being generated in the boxplot for identifying outliers; interestingly, the original *Object* variable of "Trip_Duration_in_min" can produce a box in the boxplot, which is worth further research.

In addition, based on practical significance, a maximum of over 6,000 is handled as an outlier in Google Sheets. Regarding outliers, therefore, the same treatment is done here.

```
In [9]: # Remove outliers through Boolean Masking
mask_for_non_outliers = data_deduplicated['Trip Duration (min)'] < 6000
data_without_outliers = data_deduplicated[mask_for_non_outliers]
data_without_outliers.info()

# After data cleaning, a new dataframe named "data_cleaned" is created fo
data_cleaned = data_without_outliers.reset_index(drop=True)
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 18448 entries, 0 to 20399
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Start Time                            18448 non-null  object
1   Stop Time                             18448 non-null  object
2   Start Station ID                      18448 non-null  int64
3   Start Station Name                    18448 non-null  object
4   End Station ID                        18448 non-null  int64
5   End Station Name                      18448 non-null  object
6   Bike ID                              18448 non-null  int64
7   User Type                            18448 non-null  object
8   Birth Year                           18448 non-null  int64
9   Age                                  18448 non-null  int64
10  Age Groups                           18448 non-null  object
11  Trip Duration                        18448 non-null  int64
12  Trip_Duration_in_min                 18448 non-null  object
13  Month                               18448 non-null  int64
14  Season                              18448 non-null  object
15  Temperature                          18448 non-null  int64
16  Weekday                             18448 non-null  object
17  Trip Duration (min)                  18448 non-null  int64
dtypes: int64(9), object(9)
memory usage: 2.7+ MB
```

Step 3: Perform Exploratory Data Analysis (EDA) and Data Visualizations

```
In [10]: # Gather the basic information about the cleaned data
data_cleaned.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18448 entries, 0 to 18447
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Start Time                            18448 non-null  object
1   Stop Time                             18448 non-null  object
2   Start Station ID                      18448 non-null  int64
3   Start Station Name                    18448 non-null  object
4   End Station ID                        18448 non-null  int64
5   End Station Name                      18448 non-null  object
6   Bike ID                              18448 non-null  int64
7   User Type                            18448 non-null  object
8   Birth Year                           18448 non-null  int64
9   Age                                  18448 non-null  int64
10  Age Groups                           18448 non-null  object
11  Trip Duration                        18448 non-null  int64
12  Trip_Duration_in_min                 18448 non-null  object
13  Month                               18448 non-null  int64
14  Season                              18448 non-null  object
15  Temperature                          18448 non-null  int64
16  Weekday                             18448 non-null  object
17  Trip Duration (min)                  18448 non-null  int64
dtypes: int64(9), object(9)
memory usage: 2.5+ MB
```

Question 1: What are the most popular pick-up locations across the city for NY Citi Bike rental?

```
In [11]: # Conduct Analysis of Question 1
df1 = data_cleaned['Start Station Name'].value_counts()
df1_normalized = data_cleaned['Start Station Name'].value_counts(normaliz

# Display the result of analysis
print("Analysis:")
print(f"There are {data_cleaned['Start Station Name'].nunique()} starting
print()
print("For the first question, the top 15 starting stations and their res
print()
print(df1[:15])
print('----- Proportions -----')
# Quickly grasp what proportions of the top 15 starting stations are of t
print(df1_normalized[:15])
```

Analysis:

There are 50 starting stations in total; also see the 'descriptive statist
ics' above.

For the first question, the top 15 starting stations and their respective
proportions are as follows.

Start Station Name

Grove St PATH	2319
Exchange Place	1341
Hamilton Park	1185
Sip Ave	1184
Morris Canal	768
Newport PATH	707
City Hall	622
Van Vorst Park	580
Newark Ave	554
Warren St	540
Brunswick St	525
Dixon Mills	515
Jersey & 6th St	512
Jersey & 3rd	502
Marin Light Rail	500

Name: count, dtype: int64

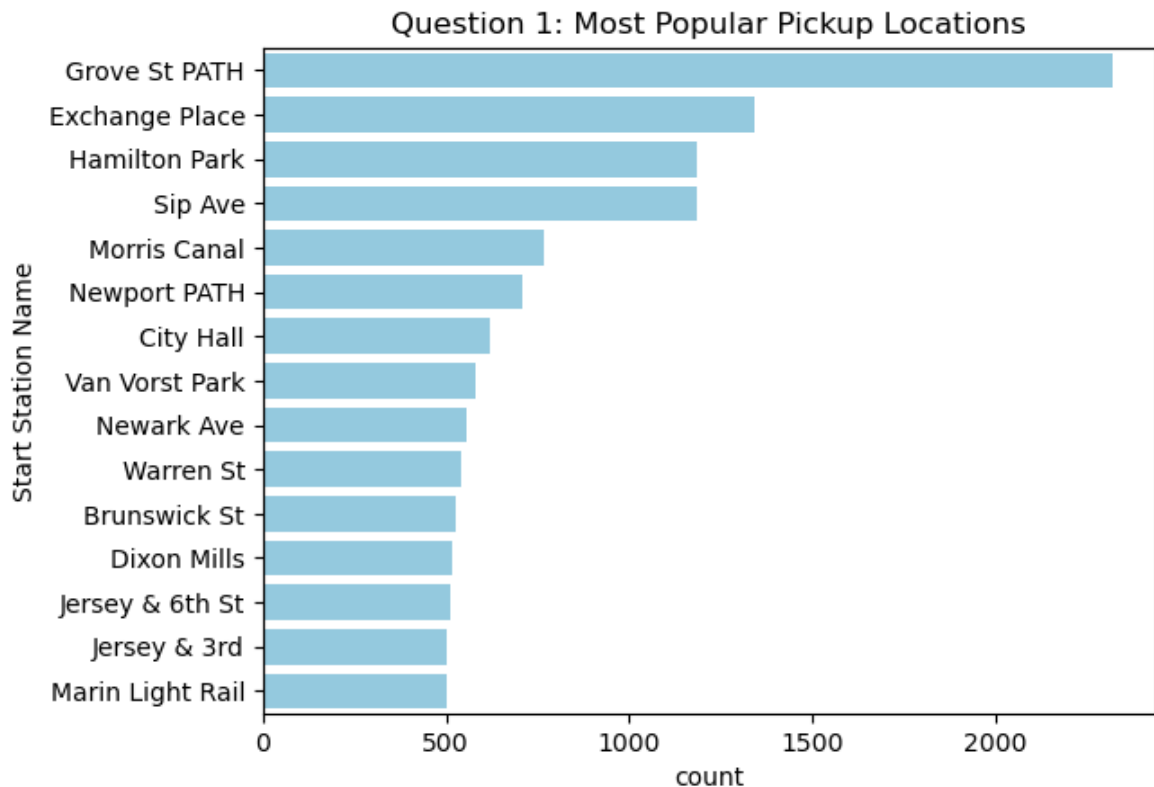
----- Proportions -----

Start Station Name

Grove St PATH	0.125705
Exchange Place	0.072691
Hamilton Park	0.064235
Sip Ave	0.064180
Morris Canal	0.041631
Newport PATH	0.038324
City Hall	0.033716
Van Vorst Park	0.031440
Newark Ave	0.030030
Warren St	0.029271
Brunswick St	0.028458
Dixon Mills	0.027916
Jersey & 6th St	0.027754
Jersey & 3rd	0.027212
Marin Light Rail	0.027103

Name: proportion, dtype: float64

```
In [12]: # Conduct Visualization for Question 1
viz1 = sns.barplot(data=df1[:15], orient="h", color='skyblue')
viz1.set_title("Question 1: Most Popular Pickup Locations")
plt.show()
```



Finding 4:

Although the de-duplicating techniques of Pandas and Google Sheets differ, the analysis results for Question 1 are the same (except for the third and fourth place rankings as well as exact numbers).

Question 2: How does the average trip duration vary across different age groups?

```
In [13]: # Conduct Analysis of Question 2
df2 = data_cleaned.groupby(['Age Groups'])['Trip Duration (min)'].mean()
df2_sorted = data_cleaned.groupby(['Age Groups'])['Trip Duration (min)'].

# Display the result of analysis
print("Analysis:")
print("For the second question, the average trip duration by the age grou")
print()
print(df2_sorted)
```


Analysis:

For the second question, the average trip duration by the age group is as follows.

Age Groups

65-74 7.350801

45-54 8.085855

25-34 9.028642

35-44 9.476480


55-64 10.200000

18-24 11.857143

75+ 47.877193

Name: Trip Duration (min), dtype: float64

```
In [14]: # Conduct Visualization for Question 2
fig, viz2 = plt.subplots(1, 2, figsize = (20,6))
viz2[0] = sns.barplot(data=df2_sorted, ax=viz2[0], color='Gold')
viz2[0].set_title("Question 2: Average Trip Duration (Ascending) by Age G
viz2[1] = sns.barplot(data=df2, ax=viz2[1], color='Gold')
viz2[1].set_title("Question 2: Average Trip Duration by Age Group (Ascend
plt.show()
```



Finding 5:

Although the de-duplicating techniques of Pandas and Google Sheets differ, the analysis results for Question 2 are the same (except for exact numbers).

Moreover, in order to quickly gain insights, an additional bar graph is generated to depict the main features when the age groups are arranged in ascending order.

Question 3: Which age group rents the most bikes?

```
In [15]: # Conduct Analysis of Question 3
df3 = data_cleaned['Age Groups'].value_counts()

# Display the result of analysis
print("Analysis:")
print("For the third question, the most active user age groups are as fol
print()
print(df3)
```

Analysis:

For the third question, the most active user age groups are as follows.

Age Groups

35-44 8376

25-34 4434

45-54 3238

55-64 1600

65-74 687


75+ 57

18-24 56

Name: count, dtype: int64

```
In [16]: ##### Conduct Visualization for Question 3
fig, viz3 = plt.subplots()
viz3_colors = sns.color_palette('pastel')
patches, texts, autotexts = viz3.pie(df3, colors=viz3_colors, autopct='%.'
viz3.legend(patches, df3.index, title="Age Groups", loc="upper right", bb
#viz3.pie(df3, labels=df3.index, colors=viz3_colors, autopct='%0.1f%', co
viz3.set_title("Question 3: Most Active User Age Groups")
plt.setp(autotexts, size=10, rotation=60)
plt.axis('equal')
plt.show()
```

Question 3: Most Active User Age Groups



Finding 6:

Although the de-duplicating techniques of Pandas and Google Sheets differ, the analysis results for Question 3 are the same (except for the exact numbers).

Meanwhile, it is worth mentioning that Google Sheets makes it much easier to create a pie chart than the *matplotlib.pyplot* and *Seaborn* libraries, considering that *Seaborn* currently does not have a direct "pieplot" function (similar to *barplot*).

Question 4: How does bike rental vary across the two user groups (one-time users vs. long-term subscribers) on different days of the week?

```
In [17]: # Conduct Analysis of Question 4
df4 = data_cleaned.groupby(['Weekday', 'User Type']).size().reset_index(name='Count') # Check if the sum of the new "Count" variable is equal to the original data
# Define the desired order of weekdays
weekday_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
df4['Weekday'] = pd.Categorical(df4['Weekday'], categories=weekday_order, ordered=True)
df4_sorted_weekday = df4.sort_values(by='Weekday').reset_index(drop=True)


# Display the result of analysis
print("Analysis:")
print("For the fourth question, the user types and their respective rental patterns are as follows.")
print()
print(df4_sorted_weekday)
```

Analysis:

For the fourth question, the user types and their respective rental patterns are as follows.

	Weekday	User Type	Count
0	Monday	One-time user	57
1	Monday	Subscriber	2715
2	Tuesday	One-time user	30
3	Tuesday	Subscriber	2647
4	Wednesday	One-time user	43
5	Wednesday	Subscriber	3590
6	Thursday	One-time user	20
7	Thursday	Subscriber	3239
8	Friday	One-time user	29
9	Friday	Subscriber	2645
10	Saturday	One-time user	102
11	Saturday	Subscriber	1612
12	Sunday	One-time user	68
13	Sunday	Subscriber	1651

```
In [18]: # Conduct Visualization for Question 4
# Seaborn's Countplot for two Categorical columns
df4_countplot = data_cleaned.loc[:, ['Weekday', 'User Type']]
df4_countplot['Weekday'] = pd.Categorical(df4_countplot['Weekday'], categories=weekday_order, ordered=True)
df4_countplot_sorted_weekday = df4_countplot.sort_values(by='Weekday').reset_index(drop=True)
viz4 = sns.countplot(x='Weekday', hue='User Type', data=df4_countplot_sorted_weekday)
viz4.set_title("Question 4: User Type & Rental Patterns")
plt.xticks(rotation=45, ha='right')
plt.show()
```



Finding 7:

Although the de-duplicating techniques of Pandas and Google Sheets differ, the analysis results for Question 4 are the same (except for the exact numbers).

Similar to the stacked stepped area chart in Google Sheets, the *Seaborn* countplot has an excellent ability to capture the primary features of two categorical variables, but in a more efficient manner.

Furthermore, regarding two categorical variables, *One-hot Encoding*, *Seaborn Catplot* (with "kind='count'", making it similar to *Seaborn Countplot*) and *Proportions.plot* (which may take a long time to produce results) are also useful techniques, especially for complex datasets.

Question 5: Does user age impact the average bike trip duration?

```
In [19]: # Conduct Analysis of Question 5
df5 = data_cleaned.loc[:, ['Age', 'Trip Duration (min)']]

# Compute the Median of the 'Age' and 'Trip Duration (min)' variables
age_median = df5['Age'].median()
trip_median = df5['Trip Duration (min)'].median()
median = pd.DataFrame({'Age': [age_median], 'Trip Duration (min)': [trip_

df5_reorganized = pd.concat([df5.describe().loc[['mean', 'min', 'max']],:]

# Display the result of analysis
```

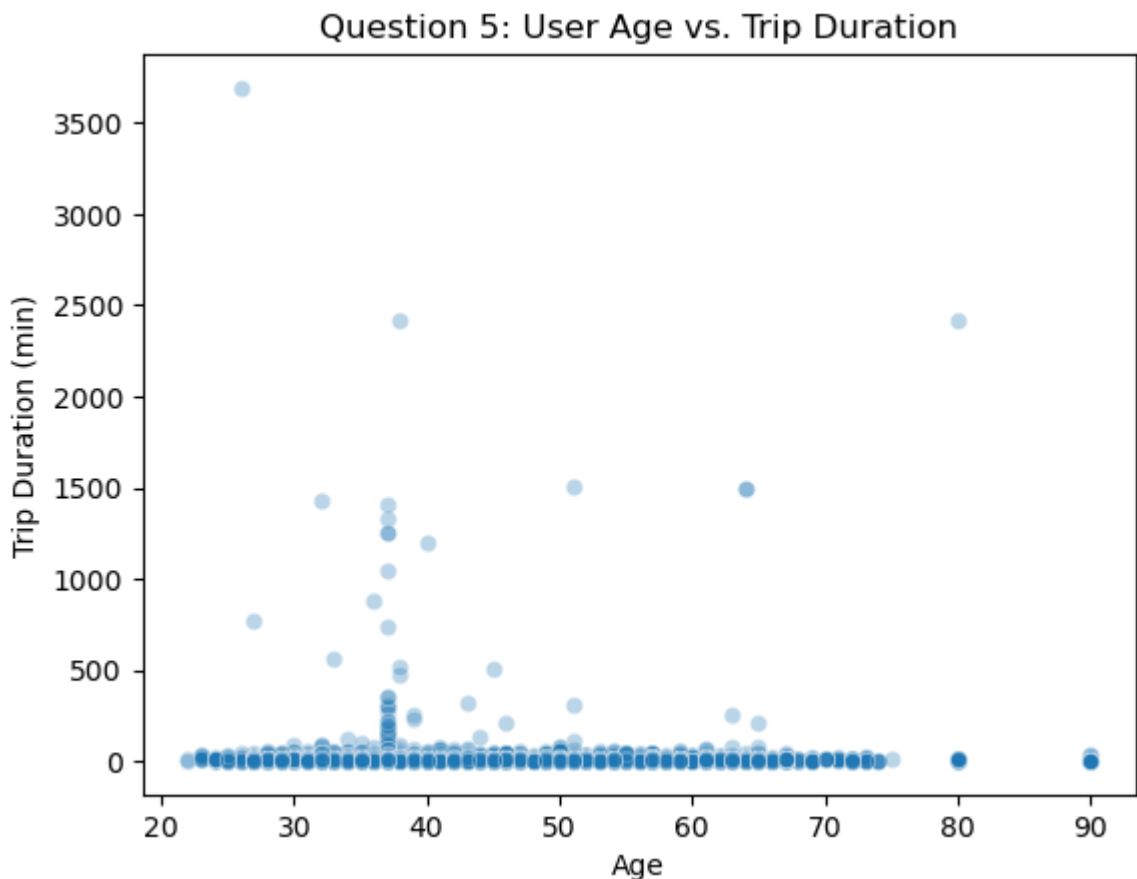
```
print("Analysis:")
print("Regarding the last question, the basic descriptive statistics of u")
print()
print(df5_reorganized)
```

Analysis:

Regarding the last question, the basic descriptive statistics of user age and trip duration are as follows.

	Age	Trip Duration (min)
mean	41.705713	9.234226
min	22.000000	1.000000
max	90.000000	3693.000000
median	39.000000	5.000000

```
In [20]: # Conduct Visualization for Question 5
viz5 = sns.scatterplot(data=df5, x='Age', y='Trip Duration (min)', alpha=
viz5.set_title("Question 5: User Age vs. Trip Duration")
plt.show()
```



Finding 8:

Although the de-duplicating techniques of Pandas and Google Sheets differ, the analysis results for Question 5 are the same (except for the exact mean values).

Moreover, the *alpha* parameter in *Seaborn Scatterplot* helps visualize many overlapping data points, thus revealing data density and improving clarity. Therefore, the same treatment is applied in Google Sheets.

Conclusions:

This project achieves an end-to-end analytical journey, beginning with a comprehensive business analysis of a bike-sharing program and culminating in a demonstration of advanced data manipulation and visualization skills.

Initially, I leveraged Google Sheets to meticulously store, clean, and analyze a dataset of over 20,400 raw records. Through exploratory data analysis using pivot tables and various visualizations, I successfully uncovered key user behavior patterns and rental trends. This analysis yielded actionable, data-driven insights that could directly help optimize station supply, target core demographics, and enhance subscriber benefits to boost weekend rentals.

To validate these findings and demonstrate my expanded technical capabilities, I recreated the entire analysis using Python. This supplementary project involved robust data cleaning, including the identification and handling of outliers, duplicates, and missing values. Utilizing Python's powerful libraries, I generated a diverse range of visualizations — from scatter plots to bar charts — which confirmed the core business findings while showcasing a professional and reproducible analytical workflow.

Ultimately, this dual-methodology project demonstrates my ability to not only translate raw data into strategic business recommendations but also to proficiently apply industry-standard analytical tools, ensuring the accuracy and integrity of every insight delivered.